





Third-person tier discrimination (neutral judge)




No self-advantage under either control




Framing-robust null + manipulation check



Generation-side amplification: discriminability tracks text source, not judge




Who is fooled by the forgery — and who sees through it?




Capability-graded calibration gradient

Overconfidence gap (95% CI)



Brier score



Behavioral prose self-authorship in the introspection debate (FOR vs AGAINST camps)

FOR privileged access

Injected-state introspection

Lindsey 2026

measured: injected internal state

Policy / behavior prediction

Naphade 2026; Binder 2024

measured: own policy / behavior

AGAINST / insufficient

Third-party-predictor bar

Song 2025

measured: compute-matched predictor standard

Reality check

Singh 2026

measured: behavioral-evidence sufficiency

★ **Behavioral prose self-authorship: NULL (this work)**

boundary on emergent introspection

measured: free-form prose authorship